

Original

Test accuracy: **0.6913**

CNN(

(features): Sequential(

(0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1))

(7): ReLU()

)



(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): Linear(in_features=1024, out_features=64, bias=True)

(2): ReLU()

(3): Linear(in_features=64, out_features=10, bias=True)



Variation 1

Test accuracy: **0.6879**

CNN(

(features): Sequential(

(0): Conv2d(3, 32, kernel_size=(5, 5), stride=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(64, 64, kernel_size=(5, 5), stride=(1, 1))

(7): ReLU()

)



(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): Linear(in_features=64, out_features=64, bias=True)

(2): ReLU()

(3): Linear(in_features=64, out_features=10, bias=True)



Variation 2

Test accuracy: **0.7146**

CNN(

(features): Sequential(

(0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(7): ReLU()

)



(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): LazyLinear(in_features=0, out_features=64, bias=True)

(2): ReLU()

(3): Linear(in_features=64, out_features=10, bias=True)



Variation 3

Test accuracy: **0.7038**

CNN(

(features): Sequential(

(0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(7): ReLU()

(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(9): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(10): ReLU()

)



(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): LazyLinear(in_features=0, out_features=64, bias=True)

(2): ReLU()

(3): Linear(in_features=64, out_features=10, bias=True)



Variation 4

Test accuracy: **0.7122**

CNN(

(features): Sequential(

(0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(7): ReLU()

(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(9): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(10): ReLU()

)



(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): LazyLinear(in_features=0, out_features=128, bias=True)

(2): ReLU()

(3): Linear(in_features=128, out_features=10, bias=True)



Variation 5

Test accuracy: **0.7026**

CNN(

(features): Sequential(

(0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(7): ReLU()

(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(9): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(10): ReLU()

)

(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)



(1): LazyLinear(in_features=0, out_features=128, bias=True)

(2): ReLU()

(3): Linear(in_features=128, out_features=64, bias=True)

(4): ReLU()

(5): Linear(in_features=64, out_features=10, bias=True)



Variation 6

Test accuracy: **0.7157**

CNN(

(features): Sequential(

(0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(7): ReLU()

(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(9): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(10): ReLU()

)

(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): LazyLinear(in_features=0, out_features=128, bias=True)

(2): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, bias=True,

track_running_stats=True)

(3): GELU(approximate='none')

(4): Dropout(p=0.5, inplace=False)

(5): Linear(in_features=128, out_features=64, bias=True)



(6): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, bias=True,

track_running_stats=True)

(7): GELU(approximate='none')

(8): Dropout(p=0.5, inplace=False)

(9): Linear(in_features=64, out_features=10, bias=True)



Variation 7

Test accuracy: **0.7512**

CNN(

(features): Sequential(

(0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(7): ReLU()

(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(9): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(10): ReLU()

)

(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): LazyLinear(in_features=0, out_features=256, bias=True)

(2): BatchNorm1d

(3): GELU(approximate='none')

(4): Dropout(p=0.5, inplace=False)

(5): Linear(in_features=256, out_features=128, bias=True)

(6): GELU(approximate='none')



(7): Dropout(p=0.5, inplace=False)

(8): Linear(in_features=128, out_features=64, bias=True)

(9): GELU(approximate='none')

(10): Dropout(p=0.5, inplace=False)

(11): Linear(in_features=64, out_features=10, bias=True)



Variation 8

Test accuracy: **0.7203**

CNN(

(features): Sequential(

(0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU()

(2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(4): ReLU()

(5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(6): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(7): ReLU()

(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(9): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(10): ReLU()

)

(classifier): Sequential(

(0): Flatten(start_dim=1, end_dim=-1)

(1): LazyLinear(in_features=0, out_features=256, bias=True)

(2): BatchNorm1d

(3): GELU(approximate='none')

(4): Dropout(p=0.5, inplace=False)

(5): Linear(in_features=256, out_features=128, bias=True)

(6): GELU(approximate='none')

(7): Dropout(p=0.5, inplace=False)

(8): Linear(in_features=128, out_features=64, bias=True)

(9): BatchNorm1d

(10): GELU(approximate='none')

(11): Dropout(p=0.5, inplace=False)

(12): Linear(in_features=64, out_features=32, bias=True)

(13): GELU(approximate='none')

(14): Dropout(p=0.5, inplace=False)

(15): Linear(in_features=32, out_features=10, bias=True)

